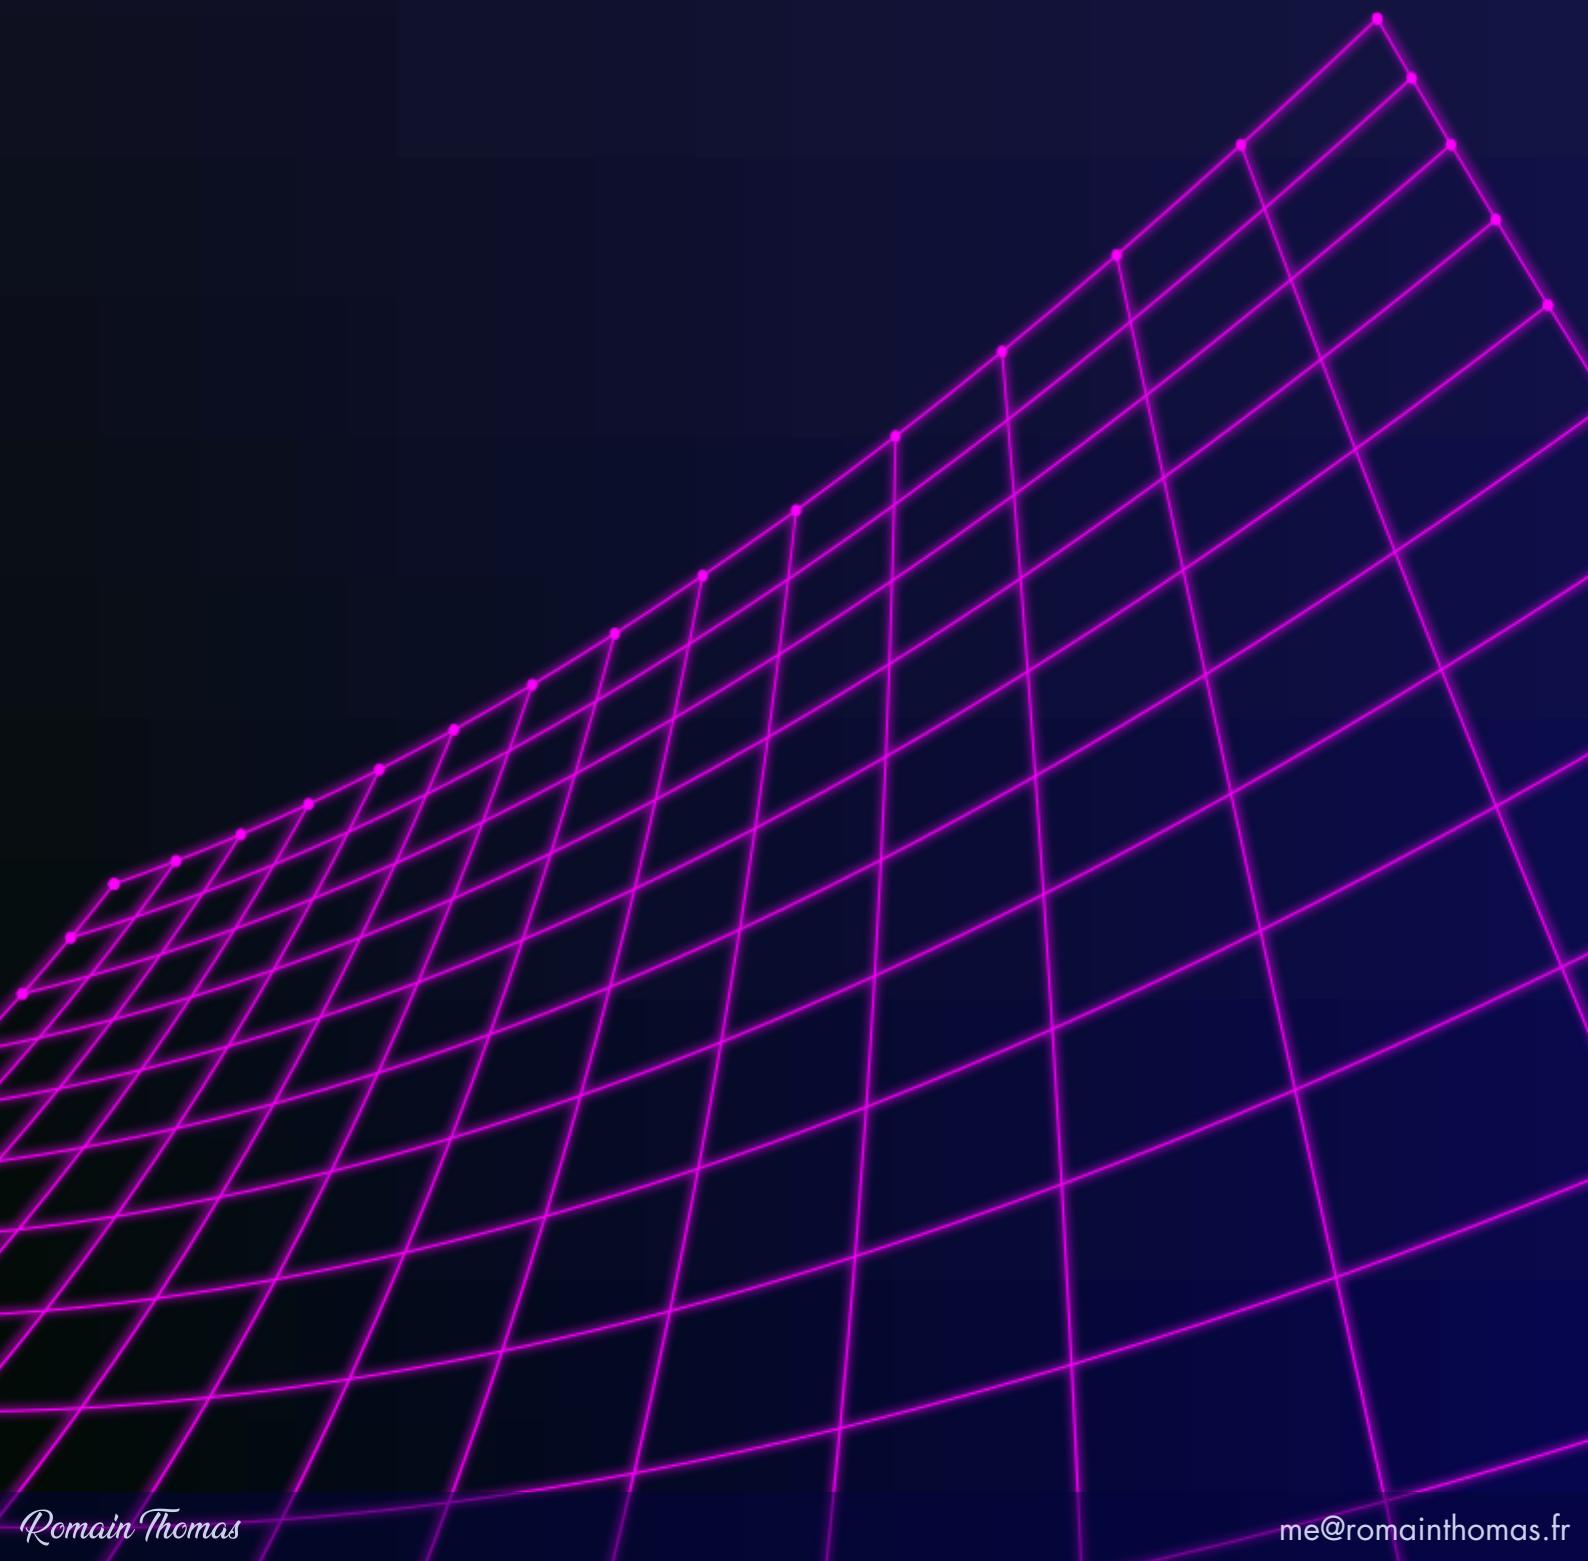


The Poor Man's Obfuscator

Pass The Salt 2022



Contents:

1	Introduction	1
2	__ unwind_info & .eh_frame	2
3	Exported Symbols	4
3.1	Creating Fake Exports Names	4
3.2	Confusing the Exports Names	5
3.3	Exports Addresses	7
4	Sections Transformations	9
4.1	ELF	9
4.2	Mach-O	11
5	Specific Transformations	14
5.1	LC_FUNCTIONS_STARTS	14
5.2	.dynsym section	16
6	Conclusion	17

1. Introduction

The purpose of this paper is to present ELF and Mach-O transformations which impact or hinder disassemblers like IDA, BinaryNinja, Ghidra, and Radare2. In particular, these transformations do not modify the assembly code or the data of the binary. The transformations are focused on modifying the executable file format structures like the sections or the symbols. As a result, the modified binaries look obfuscated as it will be shown through different examples.

All the modifications presented in this paper are based on [LIEF](#) (commit : `f8c711d`) which is a library for parsing and modifying executable file formats. The binaries referenced in the examples are based on the [Mbed TLS](#) test suite and more precisely, `programs/test/selftest.c`. This testing program is convenient to verify that the ELF/Mach-O modifications do not break the binary and do not introduce side effects.



The aim of this work is **not** to highlight or point out the limits of the disassemblers mentioned in this paper. Actually, all of them have pros and cons with different design decisions which enable us to enjoy reverse engineering at different levels.

Being not an expert on all these disassemblers, I might also miss some loading options for which the transformation would not impact the disassembler. I'll take care of updating the content of this paper based on the feedback (if any).

The content of this paper has been presented at [Pass The Salt 2022](#) on July, 5th using on the following releases of the disassemblers :

Tools	Version	Release Date
IDA	7.7.211224	January, 2022
BinaryNinja	3.0	April 2022
Radare2	5.7.0	June 2022
Ghidra	10.1.4	May 2022

2. __ unwind_info & .eh_frame

The ELF .eh_frame section and the Mach-O __ unwind_info section contain information that can be used to get a list of function addresses present in the binary.

These sections are supported and parsed by LIEF which exposes the list of functions through the `Binary.functions` attribute :

```
import lief

target = lief.parse("/bin/ls")
for function in target.functions:
    print(function.address)
```

These sections are also used by the disassemblers to get an initial *worklist* of functions where they can start the disassembling.

Thus, we can prevent the disassemblers from using the content of these sections by shuffling their content. Here is an example for the ELF sections :

```
for section in [".eh_frame", ".eh_frame_hdr"]:
    section = target.get_section(name)
    section_content = list(section.content)
    random.shuffle(section_content)
    section.content = section_content
```

IDA, BinaryNinja, and Radare2 do not seem to be impacted by this modification but Ghidra raises an exception when trying to analyse the binary :

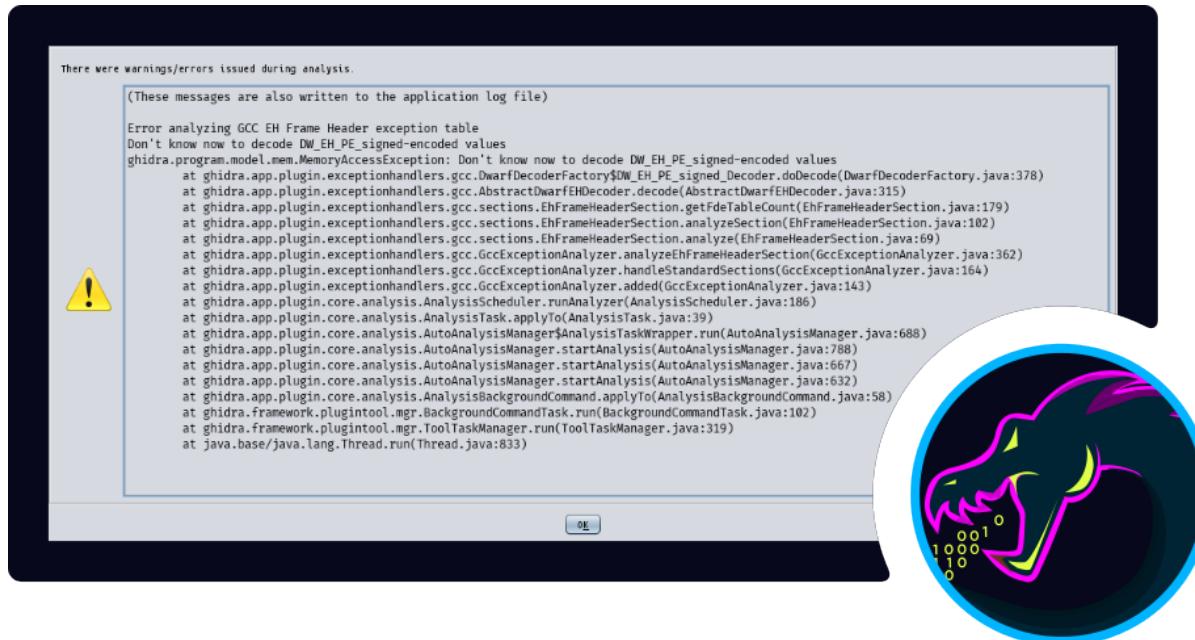


FIGURE1 – Error with Ghidra when shuffling the __unwind_info section

Tools	Impacted
IDA	No
BinaryNinja	No
Radare2	No
Ghidra	Yes

3. Exported Symbols

The first simple, efficient and universal modification we can do on executable formats is creating fake exports. Thanks to the ELF `.eh_frame` section and the Mach-O `__ unwind_info` section, we can get a – more or less – accurate list of functions’ start addresses and thus, create relevant exports.

3.1 Creating Fake Exports Names

We can start to “obfuscate” the binary by creating export names for which the name is randomized. Using this LIEF, such exports can be created by using the `add_exported_function` function.



This function is only available for the ELF and Mach-O formats and is not implemented for the PE format.

The random symbol’s name can be generated through the Python random module :

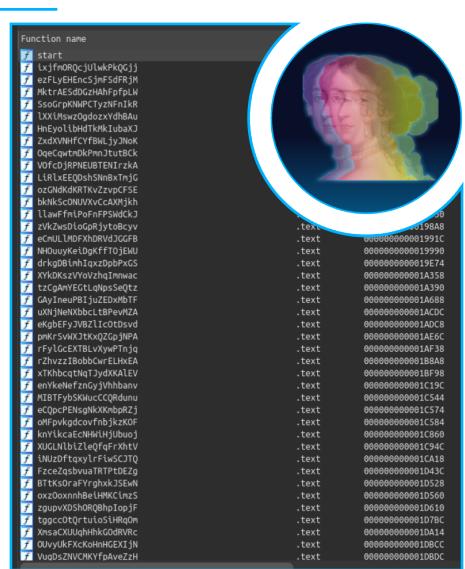
```
import lief
import random
import string

target = lief.parse("mbedtls_self_test.arm64.macho")

for function in target.functions:
    name = "".join(random.choice(string.ascii_letters) for i in range(20))
    target.add_exported_function(function.address, name)

target.write("01-mbedtls_self_test.arm64.macho")
```

In the output of IDA, BinaryNinja, etc we can observe the list of the random names as we could expect. This transformation is not very *fancy* except if the binary to protect has been compiled with a *wrong* visibility (like `-fvisibility=default`) and in which, we would like to remove or change the symbols after the compilation.



3.2 Confusing the Exports Names

Creating random export names is confusing but a reverse engineer can immediately identify that the binary has been modified or aims at being protected. In addition, it does not bring more protection compared to a regular strip of the functions.

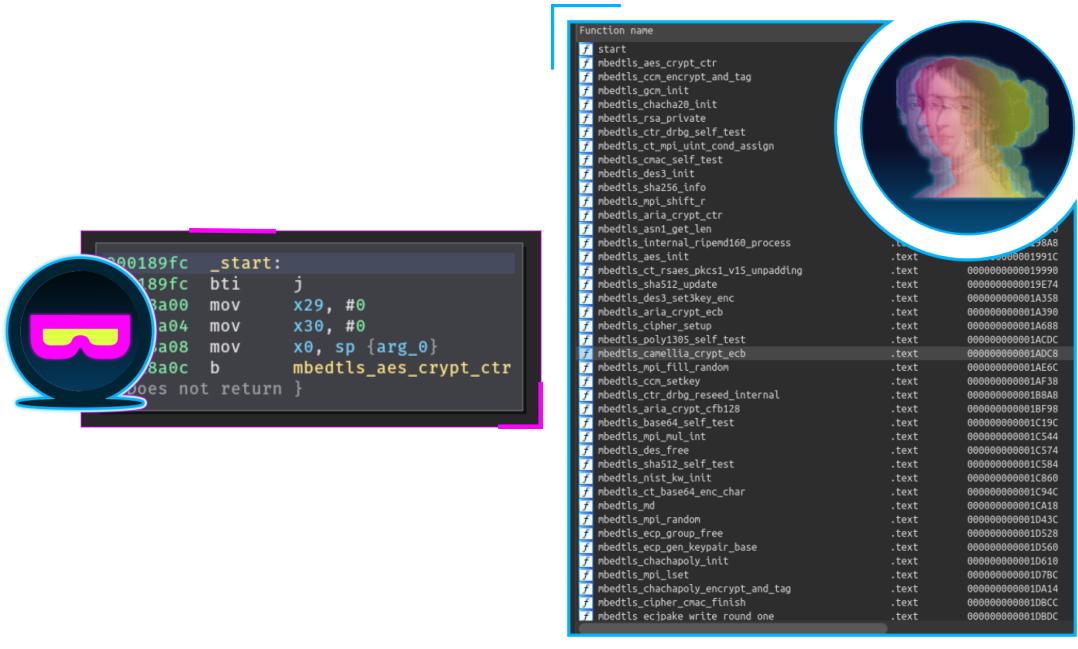
Actually, we can still leverage the exports table to add new entries but instead of using random strings, we can pick real and consistent function names from the original Mbed TLS binary.

```
target      = lief.parse("mbedtls_self_test.arm64.elf")
nonstripped = lief.parse("mbedtls_self_test.nostrip.arm64.elf")

SYMBOLS = [s.name for s in non_stripped.symbols \
           if s.name.startswith("mbedtls_")]

for function in target.functions:
    name = random.choice(SYMBOLS)
    SYMBOLS.remove(name)
    target.add_exported_function(function.address, name)
```

With such exports, it creates more confusion as it becomes difficult to distinguish the real Mbed TLS functions from the fake ones.



As we can observe in the previous figure, one function not related to `mbedtls_aes_crypt_ctr` has been renamed with this name.

To continue walking along the path of using meaningful names, we can also take names from a standard library like the `libc.so`. Compared to the `mbedtls_*` names, `libc`'s symbols are usually recognized by the disassemblers which provide type libraries for such functions. We could also be a bit more sneaky by taking C++ mangled symbols, from the LLVM libraries (for instance).

But let's continue with the `libc`'s symbols. For Android, we can collect `libc`'s symbols from the NDK and we also need to avoid exporting symbols with `libc`'s symbols already imported by the target :

```

libc = lief.parse("[ ... ]toolchains/llvm/prebuilt/linux-x86_64"
                 "/sysroot/usr/lib/aarch64-linux-android/23/libc.so")

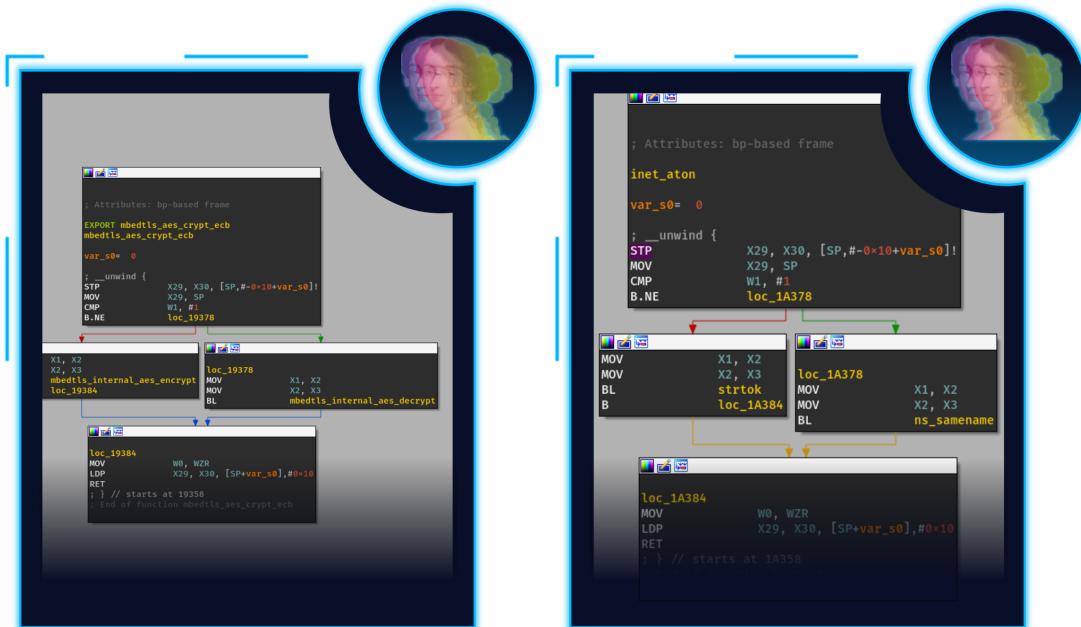
libc_symbols = {s.name for s in libc.exported_symbols}
libc_symbols -= {s.name for s in target.imported_symbols}
libc_symbols = list(libc_symbols)

for function in target.functions:
    sym = random.choice(libc_symbols)
    libc_symbols.remove(sym)

    target.add_exported_function(function.address, sym)

```

This transformation produces the following output where the figure on the left-hand side is the original unstripped function and the figure on the right-hand side, the function with the libc's symbols exported.



⚠ This transformation could be automatically defeated in different ways, like checking if the libc's symbols points in the .plt section. It could also be an interesting use case with binaries statically linked with the libc (which is forbidden on iOS and Android).

We also need to make sure that the newly exported libc's symbols are not used by the loader to actually resolve libc's functions imported by other libraries or the library itself. This can be accomplished by tweaking the symbol's visibility and the symbol's binding :

```

export = target.add_exported_function(function.address, sym)

export.binding      = lief.ELF.SYMBOL_BINDINGS.GNU_UNIQUE
export.visibility = lief.ELF.SYMBOL_VISIBILITY.INTERNAL

```

ℹ If the binary targets Linux, the binding must be set to `lief.ELF.SYMBOL_BINDINGS.WEAK` instead of `lief.ELF.SYMBOL_BINDINGS.GNU_UNIQUE`

3.3 Exports Addresses

One of the challenges when doing static analysis is to identify functions present in the binary. The binary entrypoint is obviously a good start for disassembling the code but disassemblers rely on other information from the executable file format like the exports table. It turns out that the exports table is strongly trusted by the disassemblers, whilst we can actually create entries with arbitrary symbols and arbitrary addresses.

The previous paragraphs only dealt with the export names. Since exports are always tied with an address, we can also trick this value. One of these tricks consists in creating exports with a delta value at the beginning of the function :

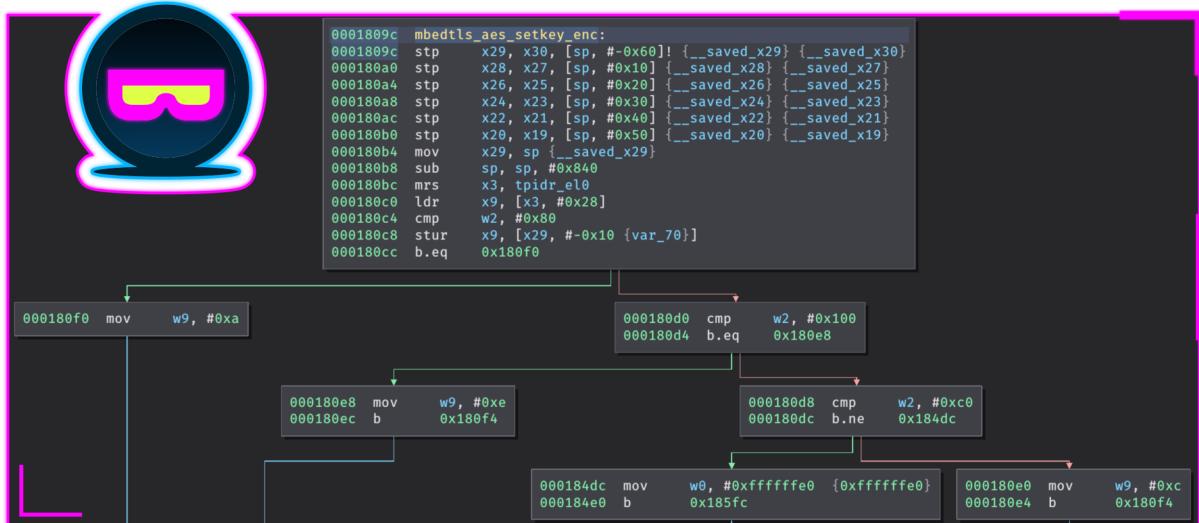
```
for function in target.functions:  
    address = function.address  
    address += random.randint(16, 32)
```

To avoid unaligned instructions, we need to make sure that the delta is aligned on four bytes :

```
for function in target.functions:  
    address = function.address  
    address += random.randint(16, 32)  
    address -= address % 4
```

With such a transformation, all the tools : IDA, BinaryNinja, Radare2, Ghidra are wrongly disassembling the code and produce an incomplete control flow graph.

Here is the impact of this transformation when opening the modified binary in BinaryNinja :



The addresses between the original binary and the modified binary are shifted by `0x1000` as a consequence of the LIEF internal mechanism to extend the exports table.

```

0001909c  sub_1909c:
0001909c  stp      x29, x30, [sp, #-0x60]! {var_60} {var_58}
000190a0  stp      x28, x27, [sp, #0x10] {var_50} {var_48}
000190a4  stp      x26, x25, [sp, #0x20] {var_40} {var_38}
000190a8  stp      x24, x23, [sp, #0x30] {var_30} {var_28}
000190ac  stp      x22, x21, [sp, #0x40] {var_20} {var_18}
000190b0  stp      x20, x19, [sp, #0x50] {var_10} {var_8}
000190b4  mov      x29, sp {var_60} {mktemp}
{ Falls through into mktemp }

```



We can observe a similar output in Radare2 :

```

[0x0001909c]> pdb
          ; XREFS: CALL 0x00019688  CALL 0x000198e4  CODE 0x00019904  CALL 0x00019958  CALL 0x0001b028  CALL 0x0001b110
          ; XREFS: CALL 0x0001b244  CALL 0x0001b424  CALL 0x0001b5a4  CALL 0x0001b734  CALL 0x0001b76c  CALL 0x0001b780
          ; XREFS: CODE 0x0002848c  CALL 0x0002a66c  CALL 0x0002a808  CALL 0x0002a9c4  CALL 0x0002ae04  CALL 0x0002b010
28: fcn.0001909c (int64_t arg1, int64_t arg2, int64_t arg3, int64_t arg_10h, int64_t arg_20h, int64_t arg_30h, int64_t
rg: 3 (vars 0, args 3)
bp: 0 (vars 0, args 0)
sp: 21 (vars 13, args 8)
    0x0001909c    fd7bbaa9    stp x29, x30, [sp, -0x60]!
    0x000190a0    fc6f01a9    stp x28, x27, [sp, 0x10]
    0x000190a4    fa6702a9    stp x26, x25, [sp, 0x20]
    0x000190a8    f85f03a9    stp x24, x23, [sp, 0x30]
    0x000190ac    f65704a9    stp x22, x21, [sp, 0x40]
    0x000190b0    f44f05a9    stp x20, x19, [sp, 0x50]
    0x000190b4    fd030091    mov x29, sp
[0x0001909c]>

```



Tools	Impacted
IDA	Yes
BinaryNinja	Yes
Radare2	Yes
Ghidra	Yes

4. Sections Transformations

The transformations based on the exports table have an important impact on the efficiency and the accuracy of the disassemblers. To enhance our *Poor Man's Obfuscator* we can also apply transformations on the sections of the ELF and Mach-O binary. Whilst both formats have sections and segments, they are used differently by the loaders. Nevertheless, we can leverage the confusion made by the disassemblers between segments and sections to mislead the output of the disassemblers.

4.1 ELF

The ELF format is likely the most error-prone format to parse compared to the PE and the Mach-O format. One of the challenges when parsing an ELF binary is the duality between sections and segments. Basically, sections are used by the compilers/linker while segments are used by the loader to run the executable¹. This means that sections should not be used to get the executable point of view of the binary. On the other hand, segments have a rough granularity over the data while sections give a better precision about the location and the meaning of the data.

For instance, it is less trivial to identify the location of the GOT based on the segments, while using the sections, the GOT is usually mirrored by the `.got` section.

We could completely get rid of the sections by removing the ELF sections table but disassemblers like IDA learned to handle such a case.



Removing the sections table could be used as an anti-debug on Linux. Indeed, gdb is not able to debug a sectionless binary.

So instead of completely removing the sections table or arbitrarily corrupting the section attributes, the idea is to swap some sections, so that we keep a certain level of consistency while still breaking the overall layout.

This swap configuration is pretty efficient against the different disassemblers :

```
SWAP_LIST = [
    ("rela.dyn", ".data.rel.ro"),
    ("got", ".got.plt"),
    ("plt", ".text"),
    ("dynsym", ".gnu.version"),
]
```

Radare2 and BinaryNinja faced difficulties with the imported symbols but the assembly code seems consistent with the original one. On the other hand, Ghidra is not able to import the binary even before launching the analysis. Lastly, IDA is able to load the binary but it corrupts the assembly code.

The swap list can be applied to the ELF binary with the following script :

1. This is not exactly true in particular on Android, where the loader enforces consistency between the `.dynamic` section and the `PT_DYNAMIC` segment.

```

target = lief.parse("mbedtls_self_test.arm64.elf")

for (lhs_name, rhs_name) in SWAP_LIST:
    print(lhs_name, rhs_name)

lhs = target.get_section(lhs_name).as_frame()
rhs = target.get_section(rhs_name).as_frame()
tmp = lhs.offset, lhs.size, lhs.name, lhs.type, lhs.virtual_address

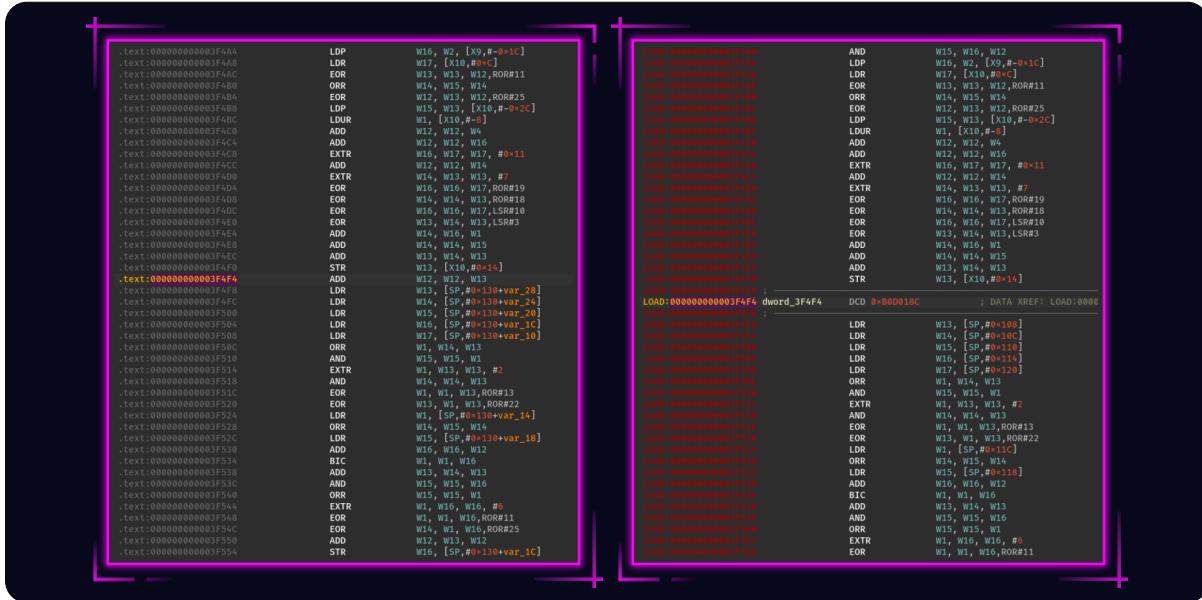
lhs.offset      = rhs.offset
lhs.size        = rhs.size
lhs.name        = rhs.name
lhs.type        = rhs.type
lhs.virtual_address = rhs.virtual_address

rhs.offset      = tmp[0]
rhs.size        = tmp[1]
rhs.name        = tmp[2]
rhs.type        = tmp[3]
rhs.virtual_address = tmp[4]

target.write("swapped_alt_mbedtls_self_test.arm64.elf")

```

The figures below show the differences in IDA between the original ELF binary and the binary with swapped sections :



This case is pretty interesting because it exists obfuscation techniques which consist in adding junk code between instructions and jumping on the real instructions with an opaque predicate. With the swapped sections, IDA is adding itself (likely because of confusion on the relocations) the junk code, making the function corrupted.



| Loading the binary with the option Use SHT disabled does not prevent the corruption.

Tools	Impacted	Note
IDA	Yes	The code is corrupted
BinaryNinja	Yes	The symbols are corrupted but the code seems consistent
Radare2	Yes	The imports are not recognized but the code seems consistent
Ghidra	Yes	We can't import the binary in Ghidra and especially, launch the analysis

4.2 Mach-O

Compared to the ELF format, the Mach-O format and its loader `dyld` enforce a stricter layout such as it is not possible to swap sections without breaking the execution of the binary. Some of the checks performed by `dyld` on the layout of the sections are defined in `dyld/dyld3/MachOAnalyze.cpp`.

Among these checks, it verifies that :

1. The section's size is not negative (or overflow)
2. The section's virtual address and virtual size is within the segment's virtual address/virtual size
3. It enforces section size alignment for special sections like `MOD_INIT_FUNC_POINTERS`

So basically, sections are stronger bound to segments than for the ELF format. Nevertheless, within the `__TEXT` segment, we can perform a small modification which consists in *virtually* shifting the beginning of the `__stubs` section over the original content of the `__text` section.



First off, these two sections are within the same `__TEXT` segment. Secondly, the transformation consists of a shift that keeps the **global** boundaries consistent with the original ones (i.e the size of all the sections does not change). The `__stubs` section is very similar to the ELF `.plt` section which is used for memoizing the resolution of the imports. In particular, the `__stubs` section contains **assembly code** (trampoline stubs) so it shares the same kind of content as the `__text` section.

Programmatically, we can perform the shift with the following piece of code :

```
SHIFT = 0x100

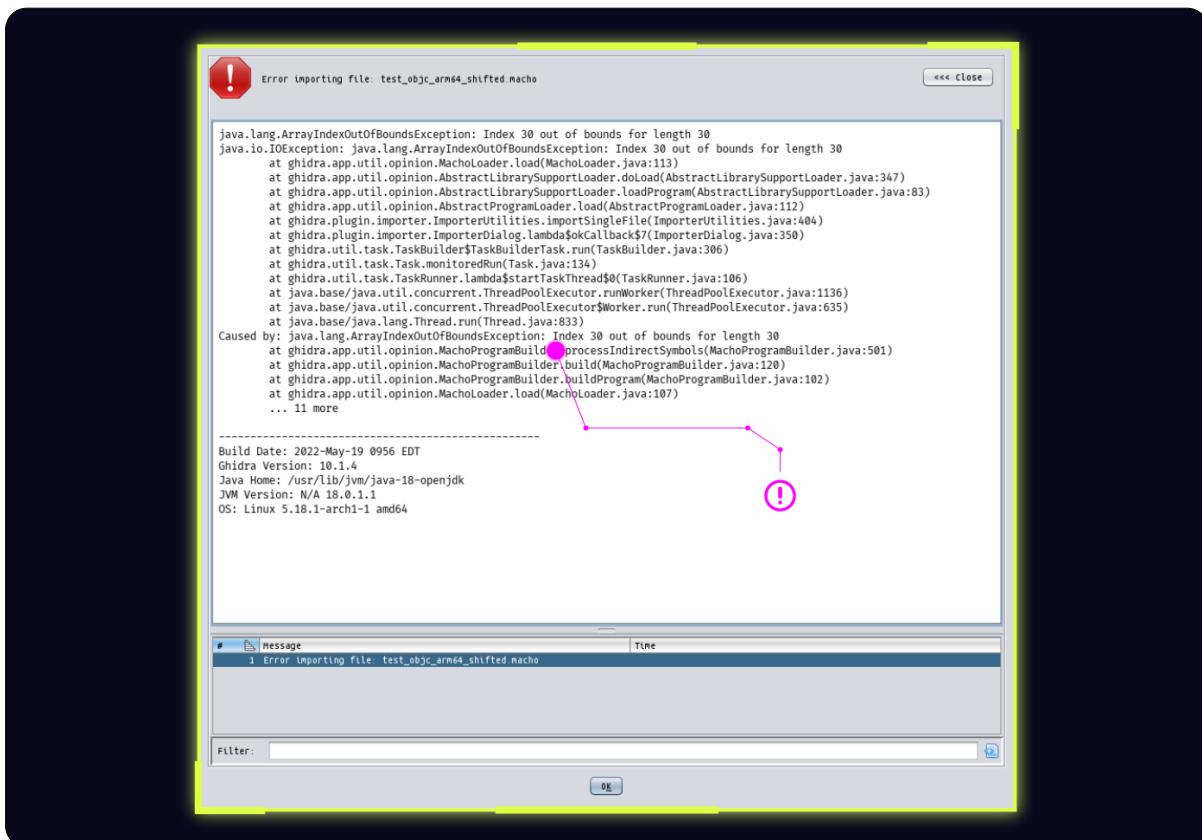
__text  = target.get_section("__text")
__stubs = target.get_section("__stubs")

# Reduce the size of the __text section
__text.size          -= SHIFT

__stubs.offset        -= SHIFT
__stubs.virtual_address -= SHIFT
__stubs.size          += SHIFT
```

Radare2 and BinaryNinja seem not impacted by this transformation while Ghidra and IDA are strongly impacted.

Ghidra simply refuses to **import** the binary. The import action comes before the analysis action so this transformation prevents Ghidra from adding a binary to a project.



On the other hand, IDA is able to load the binary but its output is confusing. In particular, the main function is broken as we can observe in the following figures :

```

; int __cdecl main(int argc, const char **argv, const char **envp)
EXPORT _main
_main

_VAR_10= -0x10

SUB      SP, SP, #0x40
STP      X29, X30, [SP,#0x40+var_10]
; End of function _main

;----- S U B R O U T I N E -----+
; void NSLog(NSSString *format, ...)
_NLog
    arg_8      = 0
    arg_20     = 0x20
    LDR      X29, X30, [SP,arg_20]
    ADD      SP, SP, #0x20 ; 0x20
;----- S U B R O U T I N E -----+
; void __cdecl _Unwind_Resume(_Unwind_Exception *exception_object)
_Unwind_Resume
    RET
;----- S U B R O U T I N E -----+
; std::allocator::allocator(void)
__ZNSt3__ allocatorC2Ev_0
    var_10    = -0x10
    var_8     = -0
    var_0     = 0
    SUB      SP, SP, #0x20
    STP      X29, X30, [SP, #0x10+var_0]
;----- S U B R O U T I N E -----+
; std::char_traits<char>::length(char const*)
__ZNSt311char_traitsifElengthPKc
    ADD      X29, SP, #0x10
    STR      XB, [SP, #0x10+var_0]
    LDR      XB, [SP, #0x10+var_0]

```

Even if the Mach-O loader enforces a stricter layout on the sections, small wisely chosen modifications enable to prevent some regular disassemblers from working correctly.

Tools	Impacted
IDA	Yes
BinaryNinja	No
Radare2	No
Ghidra	Yes

5. Specific Transformations

In the previous paragraphs, we detailed transformations based on structures shared by the two formats :

1. The exports
2. The sections

This part details other transformations that are specific to the Mach-O or the ELF format.

5.1 LC_FUNCTIONS_STARTS

The `LC_FUNCTIONS_STARTS` command is a kind of `debug` command that references a list of functions present in the binary. This command does not have an impact on `dyld` when loading the binary such as it is possible to completely corrupt its content or to modify the addresses of the functions.

Similar to the exports table, this command is used by the disassemblers to get a *trustworthy* list of functions to start disassembling.

The functions listed in this command are just a list of addresses relative to the default base address (given by the `__TEXT` segment virtual address). Using LIEF, we can apply a similar technique described in [Exports Addresses](#) by creating overlaps between two addresses :

```
functions = [f for f in LC_FUNCTION_STARTS.functions]
for idx, f in enumerate(functions):

    # Overlap 7 instructions
    if idx % 2 == 0:
        functions[idx] += 4 * 7
    else:
        functions[idx] -= 4 * 7

LC_FUNCTION_STARTS.functions = functions

bin.write("./fstart_mbedtis_self_test.arm64.macho")
```

This overlapping impacts all the disassemblers except Ghidra.

Tools	Impacted
IDA	Yes
BinaryNinja	Yes
Radare2	Yes
Ghidra	No

```
1000004b58 5f00036b cmp w2, w3
1000004b5c a1010054 b.ne 0x100004b70

1000004b40 42100034 cbz w2, 0x100004b68

1000004b44 e803022a mov w8, w2

1000004b48 09000039 ldrb w9, [x0]
1000004b4c 20000039 ldrb w10, [x1]
1000004b50 91000039 cmp w9, w10
1000004b54 e10000054 b.ne 0x100004b70

1000004b58 00040091 add x0, x0, #0x1
1000004b5c 21040091 add x1, x1, #0x1
1000004b60 00050091 subs x0, x0, #0x1
1000004b64 21ff7f54 b.ne 0x100004b48

1000004b68 00000052 mov w0, #0
1000004b6c c0035fd6 ret

1000004b70 00000012 mov w0, #xffffffff {0xffffffff}
1000004b74 c0035fd6 ret

1000004b78 int64_t sub_100004b78(int32_t* arg1, int32_t arg2, char* arg3)
1000004b78 fd7bfba9 stp x20, x30, [sp, #0x10] {__saved_x20} {__saved_x30}
1000004b7c fd030091 mov x20, sp {__saved_x20}
1000004b80 48004039 ldrb w8, [x20]
1000004b84 68000034 cbz w8, 0x100004b59

1000004b88 f7d7c188 ldp x20, x30, [sp], #0x10 {__saved_x20} {__saved_x30}
1000004b8c 39250014 bl _100000d070

1000004b90 80370004 bl _100000d030
1000004b92 00000002 mov w0, #0
1000004b98 f7d7c138 ldp x20, x30, [sp], #0x10 {__saved_x20} {__saved_x30}
1000004b9c c0035fd6 ret

1000004ba0 int64_t sub_100004ba0(int32_t arg1)
1000004ba0 8002f837 tbnz w0, #0x10, 0x100004bf0

1000004ba4 480300f0 adrp x8, 0x100006f000
1000004ba8 092543f9 ldr x9, [x8, #0x648] {data_100006f648}
1000004bb0 1000005310 adr x10, 0x100006f650
1000004bb4 1700000005 cmp x10, x11
1000004bb8 89100004 cbz x9, 0x100004be4

1000004b50 3f0010ab cmp w9, w10
1000004b54 e1000054 b.ne 0x100004b70

1000004b58 00040091 add x0, x0, #0x1
1000004b5c 21040091 add x1, x1, #0x1
1000004b60 00050091 subs x8, x8, #0x1
1000004b64 21ff7f54 b.ne 0x100004b48

1000004b68 00000052 mov w0, #0
1000004b6c c0035fd6 ret

1000004b70 00000012 mov w0, #xffffffff {0xffffffff}
1000004b74 c0035fd6 ret

1000004b78 int64_t sub_100004b78(int32_t* arg1, int32_t arg2, char* arg3, int32_t arg4)
1000004b84 68000034 cbz w8, 0x100004b90
1000004b88 fd7bc1a8 ldp x29, x30, [sp], #0x10 {arg5} {arg6}
1000004b8c 39250014 b sub_100000e070

1000004b90 a8230094 bl sub_100000d30 {sub_100004b94}
{ Falls through into sub_100004b94 }

1000004b94 int64_t sub_100004b94(int64_t arg1)
1000004b94 fd7bc1a8 ldp x29, x30, [sp], #0x10 {arg1} {arg8}
1000004b9c c0035fd6 ret

1000004b98 80 02 f8 37

1000004ba0 int64_t sub_100004ba0(int32_t arg1)
1000004ba0 480300f0 adrp x8, 0x100006f000
1000004ba4 092543f9 ldr x9, [x8, #0x648] {data_100006f648}
1000004ba8 1000005310 adr x10, 0x100006f650
1000004bb0 1700000005 nop
1000004bb4 89010004 cbz x9, 0x100004be4

1000004b50 00000002 mov x11, #0
1000004b54 4c796bb0 ldr w12, [x10, x11, lsl #0x2]
1000004b58 09100004 cmp w12, w0
1000004b5c 60010004 b.eq 0x100004bf0
```

FIGURE1 – BinaryNinja

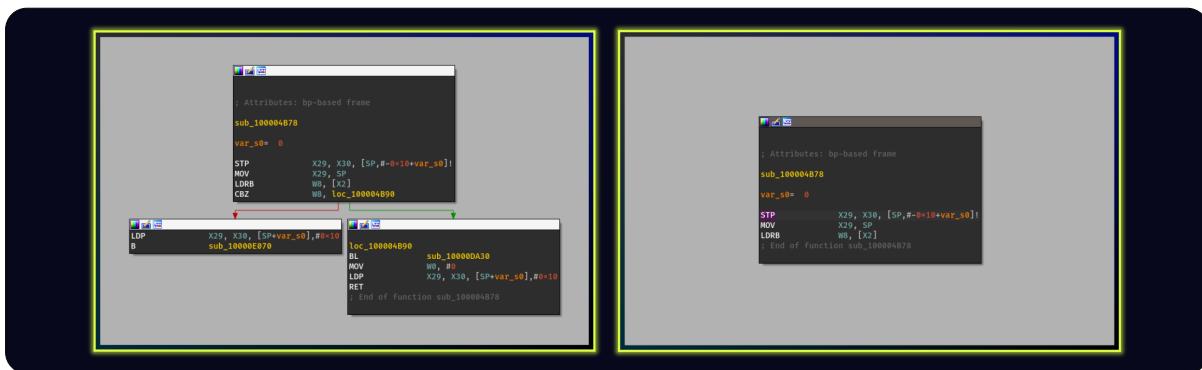


FIGURE2 – IDA

```
[0x1000004b50]> pd 20
 0x1000004b50      3f010a6b    cmp w9, w10
 0x1000004b54      e1000054    cbz w9, 0x1000004b70
 0x1000004b58      00400091    add x0, x0, 1
 0x1000004b5c      21040091    add x1, x1, 1
 0x1000004b60      08000074    subs x8, x8, 1
 0x1000004b64      21ff0054    bne w9, 0x1000004b48
 0x1000004b68      80000054    mov w0, 0
 0x1000004b72      c0035fd6    ret
 0x1000004b76      00000000    -1
 0x1000004b7a      c0035fd6    ret

12: Fcn.1000004b78 (int64_t arg1, int64_t arg2, int64_t arg3);
    ; arg int64_t arg1 @ x0
    ; arg int64_t arg2 @ x1
    ; arg int64_t arg3 @ x2
    ; var int64_t var_10h @ sp+0x0
    ; var int64_t var_14h @ sp+0x4
    ; var int64_t var_18h @ sp+0x8
 0x1000004b78      fd70bf9a    stp x29, x30, [var_10h]!
 0x1000004b7c      fd030091    mov x29, sp
 0x1000004b80      480004039   ldrb w8, [x2]
 16: sym.func.1000004b84 (int64_t arg1, int64_t arg2, int64_t arg3);
    ; arg int64_t arg1 @ x0
    ; arg int64_t arg2 @ x1
    ; arg int64_t arg3 @ x2
    ; var int64_t var_10h @ sp+0x0
    ; var int64_t var_14h @ sp+0x40
    ; var int64_t var_18h @ sp+0x70
    ; var int64_t var_30h @ sp+0x70
    ; var int64_t var_40h @ sp+0x90
 0x1000004b84      68000034    cbz w8, 0x1000004b90
 0x1000004b88      fd70c1a8    ldp x29, x30, [sp], 0x10
 0x1000004b8c      39250014    b fcn.10000e070
 0x1000004b90      a8230094    bl fcn.10000da30
 12: sym.func.1000004b94 ();
 0x1000004b94      00000000    -
 0x1000004b98      00000000    -
 0x1000004b9c      00000000    -
[0x1000004b50]> ■
(1) remain1:5.7.*
```

FIGURE3 – Radare2

5.2 .dynsym section

As already mentioned in the previous paragraphs : the ELF format is very tricky. In addition to the duality between sections and segments, counting the number of imported or the exported symbols referenced in the `.dynsym` section is not trivial.

The beginning of the symbols table associated with the imports and the exports is defined by the `DT_SYMTAB` entry which is located in the `PT_DYNAMIC` segment. The `DT_*` entries which reference a table are usually paired with another `DT_*SZ` entry which holds the size of the table. It turns out, there is an exception for the `DT_SYMTAB` entry which is not tied to another entry referencing its size.

On the other hand, the dynamic symbols table is mirrored by the `.dynsym` section which has a size. Therefore, it is appealing to use this section to count the number of entries in the table. As we mentioned in the paragraph *Sections Transformations*, ELF sections can't be trusted.

We can leverage this *feature* of the ELF format to artificially reduce the size of the `.dynsym` section :

```
dynsym      = target.get_section(".dynsym").as_frame()
sizeof      = dynsym.entry_size
osize       = dynsym.size
nsyms       = osize / sizeof
dynsym.size = sizeof * 3
```

This code artificially limits the size of the `.dynsym` section to 3 symbols.

Tools	Impacted	Note
IDA	Yes	The <code>.dynsym</code> symbols are not truncated when loading the binary with the option 'Use SHT' disabled
BinaryNinja	No	
Radare2	Yes	The <code>ia</code> command does not show all the symbols but they are correctly referenced in the assembly code (e.g. <code>reloc.puts</code> instead of <code>sym.imp.puts</code>)
Ghidra	Yes	The symbols table is truncated and it seems there is no loading option preventing it

6. Conclusion

As it has been demonstrated through this paper, file format modifications can be powerful to prevent reverse engineering tools from working correctly. File format modifications are less resilient than classical obfuscation since the original assembly code remains unchanged. On the other hand, this is a topic that is less explored than regular obfuscation and for which, it exists less tooling, automation, and literature.

You can find the different scripts used for this work in the following GitHub repository : [romainthomas/the-poor-mans-obfuscator](https://github.com/romainthomas/the-poor-mans-obfuscator)